

Assignment 9

Name: Aditya Wandhekar

Srn no: 202201449

Roll no: 34

Div:A(A2)

PS: Group and perform aggregate functions on columns in a Spark Data Frame

CODE & OUTPUT:

1) Step 1: Import Required Libraries

```
>>> from pyspark.sql import SparkSession
>>> from pyspark.sql.functions import col, sum, avg, max, min, count
```

- **SparkSession:** The entry point to Spark functionality. We use it to create a **DataFrame** and interact with the Spark cluster.
- **pyspark.sql.functions:** This module contains various built-in functions like **sum()**, **avg()**, **max()**, etc., which are used for performing operations on columns in a **DataFrame**.

2) Step 2: Create a SparkSession

```
>>> # Initialize Spark session
>>> spark = SparkSession.builder.master("local").appName("Aggregation Example").
getOrCreate()
24/11/14 10:21:22 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.
```

- **SparkSession.builder:** Starts the process of building a **SparkSession**. We specify the master node as "local" (meaning it runs on a local machine) and give the session an app name ("Aggregation Example").
- **getOrCreate():** Creates the **SparkSession** if it doesn't exist, or retrieves the existing one if it does.

3) Step 3: Create a Sample DataFrame

```
>>> data = [
...     ("Alice", "HR", 3000),
...     ("Bob", "Finance", 4000),
...     ("Alice", "HR", 3500),
...     ("Bob", "Finance", 4200),
...     ("Charlie", "IT", 5000),
...     ("Charlie", "IT", 5200)
... ]
>>> columns = ["Name", "Department", "Salary"]
>>> df = spark.createDataFrame(data, columns)
```

- **data:** A list of tuples, where each tuple contains data for a row. Each row has three values: Name, Department, and Salary.
- **columns:** A list of column names, which correspond to the data in the rows.

- **spark.createDataFrame(data, columns):** This method converts the list of tuples data into a DataFrame with the specified column names.

4) Step 4: Perform GroupBy and Aggregate Operations

```
>>> # Perform aggregation
>>> result = df.groupBy("Department").agg(
...     sum("Salary").alias("Total_Salary"),
...     avg("Salary").alias("Average_Salary"),
...     max("Salary").alias("Max_Salary"),
...     min("Salary").alias("Min_Salary"),
...     count("Salary").alias("Count")
... )
```

- **groupBy("Department"):** This groups the DataFrame by the "Department" column. It will create separate groups for each unique department value (HR, Finance, IT).
- **agg():** The aggregation function. Inside it, we define what kind of aggregation we want for each column:
 - **sum("Salary"):** Adds up the salary values for each group (department).
 - **avg("Salary"):** Calculates the average salary for each department.
 - **max("Salary"):** Finds the highest salary in each department.
 - **min("Salary"):** Finds the lowest salary in each department.
 - **count("Salary"):** Counts the number of rows (salaries) in each department.
- **alias():** Renames the results of the aggregation for better readability. For example, sum("Salary") is renamed to Total_Salary, avg("Salary") is renamed to Average_Salary, and so on.

5) Step 5: Show the Results

```
>>> result.show()
+-----+-----+-----+-----+-----+
|Department|Total_Salary|Average_Salary|Max_Salary|Min_Salary|Count|
+-----+-----+-----+-----+-----+
|      HR      |      6500   |      3250.0   |      3500   |      3000   |      2   |
|   Finance   |      8200   |      4100.0   |      4200   |      4000   |      2   |
|       IT     |     10200   |      5100.0   |      5200   |      5000   |      2   |
+-----+-----+-----+-----+-----+
```

- **show():** This command displays the resulting DataFrame, which contains the aggregated values grouped by the "Department".